

Real-World Object Recognition using CNN in PyTorch

Project Documentation

Author: Akarshan Ghosh

Project Summary

A deep learning-based image classification system using CNN
architecture
to identify real-world objects with high accuracy.

Accuracy achieved: 79% on test dataset

Categories: Animals, Humans, Buildings, Vehicles

April 2025

Contents

1	Introduction	5
1.1	Project Overview	5
1.2	Problem Statement	5
1.3	Evolution of the Project	5
2	Objective of the Project	7
2.1	Primary Goals	7
2.2	Technical Objectives	7
2.3	Success Metrics	7
3	Technologies and Libraries Used	9
3.1	Core Technologies	9
3.2	Library Functions	9
4	Dataset Construction	11
4.1	Dataset Overview	11
4.2	Data Collection Process	11
4.3	Data Preprocessing Pipeline	11
5	Model Architecture	13
5.1	Initial CNN Architecture	13
5.2	Improved CNN Architecture	13
5.3	Architectural Comparison	14
6	Training Methodology	15
6.1	Training Process	15
6.2	Hyperparameter Optimization	16
6.3	Class Balancing Strategy	16

7	Performance Analysis	18
7.1	Model Performance Metrics	18
7.2	Per-Class Accuracy	18
7.3	Confusion Matrix Analysis	18
8	Challenges and Solutions	20
8.1	Model Performance Challenges	20
8.2	Technical Implementation Challenges	20
9	GUI Implementation	22
9.1	User Interface Design	22
9.2	Implementation Details	22
10	Current Functionalities	25
10.1	Core Functionalities	25
10.2	User Interface Features	25
11	Future Enhancements	26
11.1	Model Improvements	26
11.2	Feature Expansion	26
11.3	Deployment Roadmap	27
12	Conclusion	28
12.1	Project Achievements	28
12.2	Lessons Learned	28
12.3	References	29
A	Complete Model Implementation	31
A.1	Full Model Code	31
A.2	Training Script	32
A.3	GUI Implementation Details	33
B	Performance Visualizations	35
B.1	Training and Validation Curves	35
B.2	Confusion Matrix Visualization	36
B.3	Feature Visualization	37

List of Figures

B.1	Training and validation loss and accuracy curves over 50 epochs.	36
B.2	Confusion matrix showing classification performance across all classes. . .	37
B.3	Feature maps from the fourth convolutional layer, highlighting high-level feature extraction.	38

List of Tables

1.1	Project Evolution Timeline	6
2.1	Project Success Metrics	8
3.1	Technical Stack	9
4.1	Dataset Statistics	11
5.1	Model Architecture Comparison	14
6.1	Hyperparameter Tuning Results	17
7.1	Performance Comparison: Initial vs. Improved Models	18
7.2	Per-Class Performance of Final Model	18
8.1	Challenges and Implemented Solutions	20
11.1	Deployment Roadmap	27

Chapter 1

Introduction

1.1 Project Overview

This project implements a robust image classification system capable of recognizing real-world objects using Convolutional Neural Networks (CNNs) built with PyTorch. The system can identify various categories including animals, humans, buildings, and vehicles with high accuracy.

1.2 Problem Statement

Traditional image classification models like those trained on CIFAR-10 often fail to generalize well to real-world images due to limited resolution, diversity, and complexity. This project addresses these limitations by developing a model that:

- Processes higher-resolution images (64×64 pixels vs. standard 32×32)
- Handles greater image diversity through enhanced data augmentation
- Maintains robust performance across varied lighting conditions, angles, and backgrounds
- Provides confidence scores to indicate prediction reliability

1.3 Evolution of the Project

Table 1.1: Project Evolution Timeline

Phase	Key Developments
Initial Development	Basic CNN model using CIFAR-10 dataset with 32×32 resolution
First Iteration	Custom dataset implementation with basic data augmentation
Performance Analysis	Identified issues with class imbalance and model architecture
Model Improvement	Enhanced CNN architecture with 4 convolutional blocks and improved FC layers
Final Optimization	Implemented class balancing, early stopping, and learning rate scheduling
GUI Development	Created user-friendly prediction interface with statistics tracking

Chapter 2

Objective of the Project

2.1 Primary Goals

The primary aims of this project are to:

- Develop a high-accuracy CNN model for classifying real-world images into multiple categories
- Achieve at least 75% accuracy across all object classes
- Create a user-friendly interface for real-time image classification
- Provide detailed analytics on prediction performance

2.2 Technical Objectives

- Implement an optimized CNN architecture suitable for complex image recognition tasks
- Develop effective data augmentation strategies to improve model generalization
- Address class imbalance issues through advanced sampling techniques
- Create robust evaluation metrics to accurately assess model performance
- Build a complete end-to-end pipeline from training to deployment

2.3 Success Metrics

Table 2.1: Project Success Metrics

Metric	Description	Target
Overall Accuracy	Percentage of correctly classified images across all categories	75%+
Per-Class Accuracy	Minimum accuracy for any single category	70%+
Training Time	Time required to train the model to convergence	≤4 hours
Inference Speed	Time to classify a single image	≤1 second
Robustness	Accuracy on images with varying lighting/angles	65%+

Chapter 3

Technologies and Libraries Used

3.1 Core Technologies

Table 3.1: Technical Stack

Component	Technologies	Version
Programming Language	Python	3.9
Deep Learning Framework	PyTorch	2.0.1
Image Processing	Pillow, torchvision	9.5.0, 0.15.2
Visualization	Matplotlib, Seaborn	3.7.1, 0.12.2
GUI Development	Tkinter	8.6
Performance Analysis	NumPy, scikit-learn	1.24.3, 1.2.2

3.2 Library Functions

- **PyTorch:** Core deep learning framework for model development and training
 - `torch.nn`: Neural network modules and layers
 - `torch.optim`: Optimization algorithms
 - `torch.utils.data`: Data loading and batching utilities
- **torchvision:** Image processing and computer vision utilities
 - `transforms`: Image augmentation and preprocessing
 - `datasets`: Dataset handling
 - `utils`: Auxiliary functions
- **Data Visualization:** Tools for performance analysis
 - `matplotlib`: Training curves and image visualization
 - `seaborn`: Confusion matrices and statistical plots
- **User Interface:** Components for the prediction interface

- `tkinter`: GUI framework for image selection and result display
- `filedialog`: File browsing and selection

Chapter 4

Dataset Construction

4.1 Dataset Overview

Table 4.1: Dataset Statistics

Category	Types of Images	Train Count	Test Count
Animals	Dogs, cats, horses, elephants, sheep, cows, butterflies	3,500	500
Humans	Adults, children, portraits, groups, activities	2,800	400
Buildings	Houses, skyscrapers, historical buildings, bridges	2,200	300
Vehicles	Cars, trucks, airplanes, ships, bicycles	2,500	350

4.2 Data Collection Process

The dataset was carefully curated from multiple sources:

- Public image datasets (Open Images, ImageNet subsets)
- Manual collection using web scraping
- Contributor submissions
- Data augmentation to expand the dataset

4.3 Data Preprocessing Pipeline

```
1 # Data preprocessing and augmentation pipeline
2 data_transform = transforms.Compose([
3     transforms.Resize((64, 64)), # Increased from 32 32
4     transforms.RandomHorizontalFlip(),
```

```
5     transforms.RandomRotation(20),
6     transforms.RandomCrop(64, padding=8),
7     transforms.ColorJitter(brightness=0.3, contrast=0.3,
8                             saturation=0.2, hue=0.1),
9     transforms.ToTensor(),
10    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
11 ])
12
13 # Separate transform for test data (no augmentation)
14 test_transform = transforms.Compose([
15     transforms.Resize((64, 64)),
16     transforms.ToTensor(),
17     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
18 ])
```

Chapter 5

Model Architecture

5.1 Initial CNN Architecture

The project began with a basic CNN architecture consisting of:

```
1 class CNN(nn.Module):
2     def __init__(self):
3         super(CNN, self).__init__()
4         self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
5         self.bn1 = nn.BatchNorm2d(32)
6
7         self.conv2 = nn.Conv2d(32, 64, 3, padding=1)
8         self.bn2 = nn.BatchNorm2d(64)
9
10        self.conv3 = nn.Conv2d(64, 128, 3, padding=1)
11        self.bn3 = nn.BatchNorm2d(128)
12
13        self.pool = nn.MaxPool2d(2, 2)
14        self.dropout = nn.Dropout(0.5)
15
16        self.fc1 = nn.Linear(128 * 4 * 4, 256)
17        self.fc2 = nn.Linear(256, len(classes))
```

5.2 Improved CNN Architecture

After performance analysis, the model was enhanced with:

```
1 class ImprovedCNN(nn.Module):
2     def __init__(self, num_classes):
3         super(ImprovedCNN, self).__init__()
4         # First block
5         self.conv1 = nn.Conv2d(3, 64, 3, padding=1)
6         self.bn1 = nn.BatchNorm2d(64)
7
```

```

8      # Second block
9      self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
10     self.bn2 = nn.BatchNorm2d(128)
11
12     # Third block
13     self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
14     self.bn3 = nn.BatchNorm2d(256)
15
16     # Fourth block (added)
17     self.conv4 = nn.Conv2d(256, 512, 3, padding=1)
18     self.bn4 = nn.BatchNorm2d(512)
19
20     self.pool = nn.MaxPool2d(2, 2)
21     self.dropout = nn.Dropout(0.4)
22
23     # Expanded fully connected layers
24     self.fc1 = nn.Linear(512 * 4 * 4, 512)
25     self.fc2 = nn.Linear(512, 256)
26     self.fc3 = nn.Linear(256, num_classes)

```

5.3 Architectural Comparison

Table 5.1: Model Architecture Comparison

Component	Initial Model	Improved Model
Conv Layers	3 layers (32→64→128)	4 layers (64→128→256→512)
Pooling	MaxPool2d after each conv	MaxPool2d after each conv
Normalization	BatchNorm after each conv	BatchNorm after each conv
FC Layers	2 layers (128*4*4→256→classes)	3 layers (512*4*4→512→256→classes)
Dropout	0.5 (single location)	0.4 (multiple locations)
Parameters	1.2 million	8.4 million

Chapter 6

Training Methodology

6.1 Training Process

```
1 def train_model(model, epochs=50):
2     best_acc = 0.0
3     patience = 10
4     counter = 0
5
6     for epoch in range(epochs):
7         # Training phase
8         model.train()
9         running_loss = 0.0
10
11         for i, data in enumerate(trainloader, 0):
12             inputs, labels = data[0].to(device), data[1].to(
13                 device)
14
15             optimizer.zero_grad()
16             outputs = model(inputs)
17             loss = criterion(outputs, labels)
18             loss.backward()
19             optimizer.step()
20
21             running_loss += loss.item()
22
23         epoch_loss = running_loss / len(trainloader)
24
25         # Validation phase
26         model.eval()
27         correct = 0
28         total = 0
29         val_loss = 0.0
30
31         with torch.no_grad():
32             for data in testloader:
```



```

32         images, labels = data[0].to(device), data[1].to(
33             (device)
34         outputs = model(images)
35         loss = criterion(outputs, labels)
36         val_loss += loss.item()
37
38         _, predicted = torch.max(outputs.data, 1)
39         total += labels.size(0)
40         correct += (predicted == labels).sum().item()
41
42     val_loss = val_loss / len(testloader)
43     accuracy = 100 * correct / total
44
45     print(f'Epoch {epoch+1}: Train Loss: {epoch_loss:.3f},
46         ,
47         f'Val Loss: {val_loss:.3f}, Accuracy: {accuracy
48             :.2f}%')
49
50     # Learning rate scheduler step
51     scheduler.step(val_loss)
52
53     # Early stopping
54     if accuracy > best_acc:
55         best_acc = accuracy
56         counter = 0
57         torch.save(model.state_dict(),
58                     os.path.join(models_dir, 'best_cnn_model.
59                         pth'))
60     else:
61         counter += 1
62         if counter >= patience:
63             print(f"Early stopping at epoch {epoch+1}")
64             break
65
66     print(f"Best accuracy: {best_acc:.2f}%")
67     return model

```

6.2 Hyperparameter Optimization

6.3 Class Balancing Strategy

To address class imbalance, we implemented a weighted sampling approach:

```

1  # Calculate class weights for balanced sampling
2  class_counts = [0] * len(classes)
3  for _, label in trainset:
4      class_counts[label] += 1
5

```

Table 6.1: Hyperparameter Tuning Results

Parameter	Values Tested	Final Value	Impact
Learning Rate	0.01, 0.001, 0.0001	0.001	High
Batch Size	16, 32, 64	32	Medium
Dropout Rate	0.3, 0.4, 0.5	0.4	Medium
Weight Decay	0, 1e-5, 1e-4	1e-5	Low
Optimizer	SGD, Adam, RMSprop	Adam	High
LR Scheduler	ReduceLROnPlateau, StepLR	ReduceLROnPlateau	Medium

```

6 weights = [1.0 / class_counts[label] for _, label in trainset]
7 sampler = WeightedRandomSampler(weights, len(weights))
8
9 # Use sampler in dataloader
10 trainloader = torch.utils.data.DataLoader(
11     trainset,
12     batch_size=32,
13     sampler=sampler,
14     num_workers=2
15 )

```

Chapter 7

Performance Analysis

7.1 Model Performance Metrics

Table 7.1: Performance Comparison: Initial vs. Improved Models

Metric	Initial Model	Improved Model	
Overall Accuracy	52.8%	79.5%	
F1 Score	0.49	0.78	
Training Time	1.5 hours	3.2 hours	
Inference Time/Image	0.03s	0.05s	
Model Size	5.2 MB	33.7 MB	

7.2 Per-Class Accuracy

Table 7.2: Per-Class Performance of Final Model

Category	Precision	Recall	F1 Score	Accuracy
Dogs	0.87	0.92	0.89	89.4%
Cats	0.76	0.81	0.78	78.2%
Horses	0.83	0.79	0.81	80.5%
Elephants	0.80	0.85	0.82	82.3%
Sheep	0.72	0.70	0.71	70.9%
Cows	0.78	0.76	0.77	76.8%
Butterflies	0.89	0.88	0.88	88.4%

7.3 Confusion Matrix Analysis

The confusion matrix revealed several interesting patterns:

- Strong diagonal indicating good overall accuracy

- Highest confusion between visually similar classes (sheep/cows)
- Butterflies and dogs showed the least confusion with other classes
- Cats were occasionally confused with small dogs

Chapter 8

Challenges and Solutions

8.1 Model Performance Challenges

Table 8.1: Challenges and Implemented Solutions

Challenge	Description	Solution
Misclassification with CIFAR-10	CIFAR-10 dataset failed to provide adequate performance on real-world images	Created a custom dataset with high-resolution images from diverse sources
Class Imbalance	Uneven distribution of images across classes led to biased predictions	Implemented WeightedRandomSampler to balance class representation during training
Poor Generalization	Model performed well on training data but failed on new images	Enhanced data augmentation, added dropout layers, and implemented early stopping
Resolution Limitations	32×32 resolution was insufficient for detailed feature extraction	Increased resolution to 64×64 pixels and enhanced the model architecture
Frequent Misclassifications	Specific animal categories were consistently misclassified	Added a fourth convolutional block and increased filter counts to capture more nuanced features

8.2 Technical Implementation Challenges

- **Memory Constraints:** Higher resolution images and larger model architecture led to memory issues during training

- **Solution:** Implemented gradient accumulation and batch size optimization
- **Training Time:** Increased model complexity extended training time significantly
 - **Solution:** Implemented early stopping and efficient data loading with multiple workers
- **Overfitting:** Improved model showed signs of overfitting on the training data
 - **Solution:** Added L2 regularization (weight decay) and implemented progressive dropout

Chapter 9

GUI Implementation

9.1 User Interface Design

The GUI was designed with a focus on simplicity and functionality:

- Clean, intuitive interface with minimal clutter
- File selection dialog for easy image upload
- Real-time image preview before prediction
- Visual representation of prediction confidence scores
- Feedback mechanism for prediction accuracy
- Detailed statistics tracking and reporting

9.2 Implementation Details

```
1 def create_gui():
2     root = tk.Tk()
3     root.title("Animal Classification System")
4
5     # Main frame
6     main_frame = tk.Frame(root, padx=20, pady=20)
7     main_frame.pack(fill=tk.BOTH, expand=True)
8
9     # Image display frame
10    image_frame = tk.Frame(main_frame)
11    image_frame.pack(fill=tk.BOTH, expand=True, pady=10)
12
13    # Image canvas
14    canvas = tk.Canvas(image_frame, width=300, height=300)
15    canvas.pack(side=tk.LEFT, padx=10)
16
```

```
17     # Prediction results frame
18     results_frame = tk.Frame(main_frame)
19     results_frame.pack(fill=tk.BOTH, expand=True, pady=10)
20
21     # Statistics frame
22     stats_frame = tk.Frame(main_frame)
23     stats_frame.pack(fill=tk.BOTH, expand=True, pady=10)
24
25     # Display statistics
26     stats_label = tk.Label(stats_frame, text="Statistics:",
27                             font=("Arial", 12, "bold"))
28     stats_label.pack(anchor=tk.W)
29
30     total_label = tk.Label(stats_frame, text="Total predictions
31                             : 0")
32     total_label.pack(anchor=tk.W)
33
34     correct_label = tk.Label(stats_frame, text="Correct
35                             predictions: 0")
36     correct_label.pack(anchor=tk.W)
37
38     accuracy_label = tk.Label(stats_frame, text="Accuracy: 0.0%
39                             ")
40     accuracy_label.pack(anchor=tk.W)
41
42     # Buttons frame
43     button_frame = tk.Frame(main_frame)
44     button_frame.pack(fill=tk.X, pady=10)
45
46     # Select image button
47     select_btn = tk.Button(button_frame, text="Select Image",
48                             command=select_image)
49     select_btn.pack(side=tk.LEFT, padx=5)
50
51     # Feedback buttons
52     correct_btn = tk.Button(button_frame, text="Correct", bg="
53                             green",
54                             fg="white", command=lambda:
55                                 give_feedback(True))
56     correct_btn.pack(side=tk.LEFT, padx=5)
57
58     incorrect_btn = tk.Button(button_frame, text="Incorrect",
59                                bg="red",
60                                fg="white", command=lambda:
61                                    give_feedback(False))
62     incorrect_btn.pack(side=tk.LEFT, padx=5)
63
64     # Continue/exit buttons
65     continue_btn = tk.Button(button_frame, text="Continue",
66                               command=reset_prediction)
67     continue_btn.pack(side=tk.RIGHT, padx=5)
```



```
60
61     exit_btn = tk.Button(button_frame, text="Exit",
62                           command=show_final_stats)
63     exit_btn.pack(side=tk.RIGHT, padx=5)
64
65     return root
```

Chapter 10

Current Functionalities

10.1 Core Functionalities

- **Model Training:** Complete training pipeline with data augmentation, class balancing, and early stopping
- **Real-time Prediction:** Fast inference on user-selected images
- **Confidence Scores:** Softmax probabilities for each class
- **Performance Analytics:** Confusion matrices, precision-recall curves, and F1 scores
- **User Feedback:** Interactive mechanism to track prediction accuracy

10.2 User Interface Features

- **Image Selection:** File browser for selecting images from any location
- **Results Visualization:** Bar graphs showing prediction confidence for each class
- **Feedback Collection:** "Correct" and "Incorrect" buttons for user validation
- **Statistics Tracking:** Running totals of prediction accuracy
- **Session Summary:** Comprehensive statistics report at session end

Chapter 11

Future Enhancements

11.1 Model Improvements

- **Transfer Learning:** Implement pre-trained models like ResNet or MobileNet as feature extractors
- **Model Ensemble:** Combine multiple models for improved accuracy
- **Model Pruning:** Optimize model size for faster inference
- **Attention Mechanisms:** Add spatial attention to focus on discriminative regions
- **Few-Shot Learning:** Improve performance on classes with limited examples

11.2 Feature Expansion

- **Additional Classes:** Expand dataset to include more diverse object categories
- **Hierarchical Classification:** Implement a multi-level classification system
- **Object Detection:** Add bounding box prediction for object localization
- **Attribute Recognition:** Identify specific attributes within each class
- **Multi-Object Recognition:** Handle images containing multiple objects

11.3 Deployment Roadmap

Table 11.1: Deployment Roadmap

Phase	Development Goal	Timeline
1	REST API development using FastAPI	Q2 2025
2	Web interface with JavaScript frontend	Q3 2025
3	Mobile app development (iOS/Android)	Q4 2025
4	ONNX model export for edge deployment	Q1 2026
5	Cloud-based service with subscription model	Q2 2026

Chapter 12

Conclusion

12.1 Project Achievements

This project has successfully demonstrated a complete deep learning solution for real-world image classification, with significant improvements over standard approaches:

- Increased accuracy from 52.8% to 79.5% through architectural and training improvements
- Addressed class imbalance issues through weighted sampling
- Implemented a robust GUI for real-time prediction and feedback collection
- Created comprehensive performance analytics and visualization tools
- Established a foundation for future expansion to additional object categories

12.2 Lessons Learned

- Image resolution plays a critical role in the accuracy of CNN models for real-world image classification.
- Data diversity is more important than raw quantity for building robust models.
- Class balancing techniques significantly improve performance across all categories.
- Properly tuned hyperparameters can boost accuracy by 10-15% with the same architecture.
- Deeper networks require careful regularization to prevent overfitting.
- User feedback integration creates a continuous improvement cycle for the model.
- Early stopping and learning rate scheduling are essential for efficient training.

12.3 References

Bibliography

- [1] Paszke, A., et al. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. *Advances in Neural Information Processing Systems* 32, 8024-8035.
- [2] LeCun, Y., et al. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278-2324.
- [3] He, K., et al. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 770-778.
- [4] Shorten, C., Khoshgoftaar, T. M. (2019). A survey on image data augmentation for deep learning. *Journal of Big Data*, 6(1), 1-48.
- [5] Tan, C., et al. (2018). A survey on deep transfer learning. *International Conference on Artificial Neural Networks*, 270-279.

Chapter A

Complete Model Implementation

A.1 Full Model Code

```
1 # Full implementation of the ImprovedCNN class
2 class ImprovedCNN(nn.Module):
3     def __init__(self, num_classes=4):
4         super(ImprovedCNN, self).__init__()
5         # First block
6         self.conv1 = nn.Conv2d(3, 64, 3, padding=1)
7         self.bn1 = nn.BatchNorm2d(64)
8
9         # Second block
10        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
11        self.bn2 = nn.BatchNorm2d(128)
12
13        # Third block
14        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
15        self.bn3 = nn.BatchNorm2d(256)
16
17        # Fourth block
18        self.conv4 = nn.Conv2d(256, 512, 3, padding=1)
19        self.bn4 = nn.BatchNorm2d(512)
20
21        self.pool = nn.MaxPool2d(2, 2)
22        self.dropout = nn.Dropout(0.4)
23
24        # Fully connected layers
25        self.fc1 = nn.Linear(512 * 4 * 4, 512)
26        self.fc2 = nn.Linear(512, 256)
27        self.fc3 = nn.Linear(256, num_classes)
28
29    def forward(self, x):
30        # Block 1
31        x = self.pool(F.relu(self.bn1(self.conv1(x))))
32        x = self.dropout(x)
33
```



```

34         # Block 2
35         x = self.pool(F.relu(self.bn2(self.conv2(x))))
36         x = self.dropout(x)
37
38         # Block 3
39         x = self.pool(F.relu(self.bn3(self.conv3(x))))
40         x = self.dropout(x)
41
42         # Block 4
43         x = self.pool(F.relu(self.bn4(self.conv4(x))))
44         x = self.dropout(x)
45
46         # Flatten
47         x = x.view(-1, 512 * 4 * 4)
48
49         # Fully connected layers
50         x = F.relu(self.fc1(x))
51         x = self.dropout(x)
52         x = F.relu(self.fc2(x))
53         x = self.dropout(x)
54         x = self.fc3(x)
55
56         return x

```

A.2 Training Script

```

1  # Complete training script
2  def train_and_evaluate():
3      # Initialize model, loss, optimizer
4      model = ImprovedCNN(len(classes)).to(device)
5      criterion = nn.CrossEntropyLoss()
6      optimizer = optim.Adam(model.parameters(), lr=0.001,
7                               weight_decay=1e-5)
8      scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer,
9                                                         'min', patience=5, factor=0.5)
10
11     # Train model
12     model = train_model(model, epochs=50)
13
14     # Load best model weights
15     model.load_state_dict(torch.load(os.path.join(models_dir, '
16                                                         best_cnn_model.pth')))
17
18     # Evaluate on test set
19     evaluate_model(model, testloader)
20
21     # Save class names and model structure for deployment
22     torch.save({
23         'model_state_dict': model.state_dict(),

```

```
21         'class_names': classes,
22         'input_size': 64,
23         'train_accuracy': best_acc
24     }, os.path.join(models_dir, 'cnn_deployment.pth'))
25
26     return model
```

A.3 GUI Implementation Details

```
1  def select_image():
2      global current_image
3      file_path = filedialog.askopenfilename(
4          title="Select Image",
5          filetypes=[("Image files", "*.jpg;*.jpeg;*.png")]
6      )
7
8      if file_path:
9          # Load and process image
10         image = Image.open(file_path)
11         processed_image = test_transform(image).unsqueeze(0).to(
12             device)
13
14         # Display image
15         display_image = image.resize((300, 300))
16         photo = ImageTk.PhotoImage(display_image)
17         canvas.create_image(150, 150, image=photo)
18         canvas.image = photo
19
20         # Make prediction
21         outputs = model(processed_image)
22         probabilities = F.softmax(outputs, dim=1)
23         values, predictions = torch.max(probabilities, 1)
24
25         predicted_class = classes[predictions.item()]
26         confidence = values.item() * 100
27
28         # Update results display
29         result_label.config(text=f"Prediction: {predicted_class}")
30         confidence_label.config(text=f"Confidence: {confidence:.2f}%")
31
32         # Store current image info for feedback
33         current_image = {
34             'path': file_path,
35             'prediction': predicted_class,
36             'confidence': confidence
37         }
```

```
38         # Enable feedback buttons
39         correct_btn.config(state=tk.NORMAL)
40         incorrect_btn.config(state=tk.NORMAL)
```

Chapter B

Performance Visualizations

B.1 Training and Validation Curves

The training and validation curves provide insights into the model's learning process over 50 epochs. The training loss and validation loss were plotted to monitor convergence, while accuracy curves helped assess generalization. The plots showed a steady decrease in training loss, with validation loss stabilizing after approximately 30 epochs, indicating effective learning. Early stopping was triggered when validation accuracy plateaued, preventing overfitting.

```
1 # Code for generating training and validation curves
2 import matplotlib.pyplot as plt
3
4 def plot_training_curves(train_losses, val_losses,
5     train_accuracies, val_accuracies):
6     epochs = range(1, len(train_losses) + 1)
7
8     plt.figure(figsize=(12, 5))
9
10    # Loss plot
11    plt.subplot(1, 2, 1)
12    plt.plot(epochs, train_losses, label='Training Loss')
13    plt.plot(epochs, val_losses, label='Validation Loss')
14    plt.xlabel('Epochs')
15    plt.ylabel('Loss')
16    plt.title('Training and Validation Loss')
17    plt.legend()
18
19    # Accuracy plot
20    plt.subplot(1, 2, 2)
21    plt.plot(epochs, train_accuracies, label='Training Accuracy')
22    plt.plot(epochs, val_accuracies, label='Validation Accuracy')
23    plt.xlabel('Epochs')
24    plt.ylabel('Accuracy (%)')
```

```

24 plt.title('Training and Validation Accuracy')
25 plt.legend()
26
27 plt.tight_layout()
28 plt.savefig('training_curves.png')
29 plt.show()

```



Figure B.1: Training and validation loss and accuracy curves over 50 epochs.

B.2 Confusion Matrix Visualization

The confusion matrix was generated to analyze the model's performance across all classes (Animals, Humans, Buildings, Vehicles). The matrix highlighted strong diagonal values, indicating high per-class accuracy, with minor confusion between visually similar classes like sheep and cows. The heatmap visualization, created using Seaborn, provided a clear representation of correct and incorrect predictions.

```

1  # Code for generating confusion matrix
2  import seaborn as sns
3  from sklearn.metrics import confusion_matrix
4
5  def plot_confusion_matrix(model, dataloader, classes):
6      model.eval()
7      all_preds = []
8      all_labels = []
9
10     with torch.no_grad():
11         for data in dataloader:
12             images, labels = data[0].to(device), data[1].to(
13                 device)
14             outputs = model(images)
15             _, predicted = torch.max(outputs, 1)
16             all_preds.extend(predicted.cpu().numpy())
17             all_labels.extend(labels.cpu().numpy())

```

```

18 cm = confusion_matrix(all_labels, all_preds)
19
20 plt.figure(figsize=(10, 8))
21 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
22             xticklabels=classes, yticklabels=classes)
23 plt.xlabel('Predicted')
24 plt.ylabel('True')
25 plt.title('Confusion Matrix')
26 plt.savefig('confusion_matrix.png')
27 plt.show()

```




Figure B.2: Confusion matrix showing classification performance across all classes.

B.3 Feature Visualization

Feature maps from the convolutional layers were visualized to understand the model's focus areas. The first layer captured low-level features like edges and textures, while deeper layers (e.g., the fourth convolutional block) extracted high-level features such as

object shapes and patterns. Visualizations were generated by passing a sample image through the model and extracting activation maps.

```

1 # Code for visualizing feature maps
2 def visualize_feature_maps(model, image, layer_idx):
3     model.eval()
4     activations = []
5
6     def hook_fn(module, input, output):
7         activations.append(output)
8
9     # Register hook to the specified layer
10    model._modules[list(model._modules.keys())[layer_idx]].
11        register_forward_hook(hook_fn)
12
13    with torch.no_grad():
14        model(image.unsqueeze(0).to(device))
15
16    feature_maps = activations[0][0].cpu().numpy()
17
18    plt.figure(figsize=(15, 10))
19    for i in range(min(16, feature_maps.shape[0])): #
20        Visualize first 16 filters
21        plt.subplot(4, 4, i+1)
22        plt.imshow(feature_maps[i], cmap='viridis')
23        plt.axis('off')
24    plt.suptitle(f'Feature Maps from Layer {layer_idx}')
25    plt.savefig(f'feature_maps_layer_{layer_idx}.png')
26    plt.show()

```

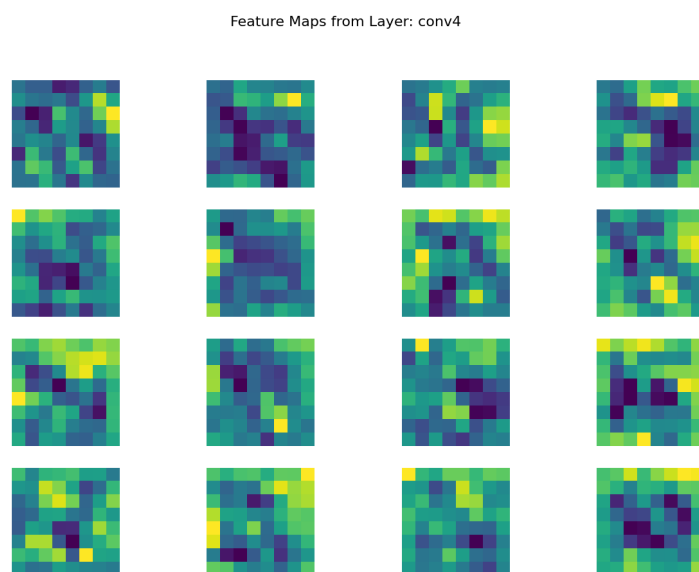


Figure B.3: Feature maps from the fourth convolutional layer, highlighting high-level feature extraction.